

Laboratory Assessment 1: Parallel Computin

Student 1076589

August 7, 2017

1 Accessing multi-dimensional arrays

By compiling and running ¹ lab11.cpp, the following table and graph were observed.

Accessing multi-dimensional arrays		
Array size	row-wise access	column-wise access
100.000000	0.000040	0.000043
200.000000	0.000160	0.000165
300.000000	0.000347	0.000429
400.000000	0.000592	0.000726
500.000000	0.000886	0.001591
600.000000	0.001183	0.002614
700.000000	0.001634	0.003891
800.000000	0.002100	0.005086
900.000000	0.002677	0.006899
1000.000000	0.003307	0.009133
1100.000000	0.004196	0.014782
1200.000000	0.005072	0.013549
1300.000000	0.006095	0.016195
1400.000000	0.007211	0.018396
1500.000000	0.008448	0.032558
1600.000000	0.009751	0.024737
1700.000000	0.011200	0.028441
1800.000000	0.012652	0.031516
1900.000000	0.014356	0.041240
2000.000000	0.016065	0.039652

Figure 1: Table of array dimension vs access time by row-wise and column-wise access

¹using -std=c++11 flag




Figure 2: Line plot of array dimension vs access time, comparing row-wise and column-wise access

It can be seen that for different array dimensions, accessing/setting data row-wise generally performs better compared to column-wise.

This performance difference can be explained by the way that multi-dimensional arrays are stored and by how instructions are executed in von Neumann model machines

In memory, a two-dimensional array is represented as a block of addresses pointing to the blocks of actual data. This implies that, although each row is stored in a continuous block of memory, the entire array may not be.

This implies that accessing multi-dimensional row-wise takes advantage of sequential locality², improving performance over column-wise access.

²spatial locality when data elements are arranged and accessed linearly

2 Display output from threads

After compiling lab12.cpp and running it multiple times, it was noted that the output varied with each execution. The order in which each thread executed was unpredictable and one thread would occasionally execute inbetween another.

This is because a threads execution is controlled by the scheduler. Based on the state of the machine, the scheduler will "decide" which thread to let execute. This causes a generally unpredictable output. At the same time, since the instructions that the threads are executing are not atomic, the scheduler may decide to interrupt a thread and allow another to execute, which may produce "inconsistent" output.